

NOTES ON PCPS AND UNIQUE GAMES

EDWARD KIM

ABSTRACT. We expound on the basic theory of approximating NP-hard problems. Starting with the PCP Theorem, we introduce background for the Unique Games Conjecture (UGC) and the optimality of the Goemans-Williamson Algorithm for MAXCUT assuming the conjecture is true.

CONTENTS

1. Approximation Algorithms and CSPs	2
1.1. Definitions and Examples	2
1.2. Constraint Satisfaction Problems	3
1.3. Gap Problems	4
2. The PCP Theorem	5
2.1. Intuitions	5
2.2. Equivalence of PCP Theorems	7
3. Label-Cover and Projection Games	8
3.1. Raz's Parallel Repetition Theorem	9
4. Håstad's 3-bit PCP	11
4.1. Inapproximability Results for MaxE3Lin and Max3Sat	11
4.2. The Long Code	12
4.3. Aside on Dictatorship Testing	12
4.4. The Code Consistency Test	14
4.5. Folded Functions	15
4.6. Soundness and Decoding a Labeling	15
5. Unique Games	17
5.1. Definitions	17
6. UG-hardness of MAXCUT	19
6.1. Intuitions	19
6.2. Goemans-Williamson Algorithm for MaxCut	19
6.3. MaxCut is UG-hard	20
6.4. Dictatorship Testing and Majority is the Stablest (MIS)	21
6.5. The 2-query PCP	23
References	26

These notes borrow heavily from Arora and Barak's textbook [1], O'Donnell's textbook [15], Prahladh Harsha's lecture notes on PCPs and Unique Games given at DIMACS and TIFR[8], [9], and Venkatesan Guruswami and Ryan O'Donnell's Course at UWash [7] . Supplementary reading can be found in Luca Trevisan's exposition on the UGC [17] and Subhash Khot's survey [13].

1. APPROXIMATION ALGORITHMS AND CSPs

1.1. Definitions and Examples. The theory of approximating NP-hard problems roots itself in the following question: “Is it possible to efficiently approximate NP-complete problem to some arbitrary degree of accuracy?” Since the Cook-Levin Theorem demonstrated that the SAT decision problem is NP-complete, the question could be rephrased as “Can we find an efficient algorithm for SAT which, given an input formula, outputs an assignment which satisfies a $1 - \delta$ fraction of the clauses for any constant $\delta > 0$?” Towards this end, let us first introduce a few motivating examples which will find utility in our forthcoming analysis.

Consider Max3CNF, the problem of, given a 3CNF formula φ as input, outputting an assignment which maximizes the number of clauses satisfied in φ . For any assignment x to the variables, say as an n -bit input string while viewing $\varphi : \{0, 1\}^n \rightarrow \{0, 1\}$ as a boolean function, there is some $0 \leq \rho_{\varphi, x} \leq 1$ representing the fraction of clauses satisfied in φ . Take $\text{Opt}(\varphi)$ to be the maximum value over all such assignments:

$$\text{Opt}(\varphi) = \max_{x \in \{0, 1\}^n} \rho_{\varphi, x}$$

A polynomial-time algorithm which solves Max3CNF is one which takes a 3CNF φ as input and outputs $\text{Opt}(\varphi)$. Naturally, Max3CNF is NP-hard since its corresponding decision problem 3CNF is NP-complete (φ is satisfiable iff $\text{Opt}(\varphi) = 1$). However, we can ask if there exists an algorithm A which outputs an assignment which satisfies at least some $\beta \cdot \text{Opt}(\varphi)$ fraction of the clauses of φ for some $\beta \leq 1$. To formalize this:

Definition 1.1. For $\beta \leq 1$, a polynomial-time algorithm A is deemed as a β -approximation algorithm for Max3CNF if $A(\varphi)$ outputs an assignment which satisfies at least $\beta \cdot \text{Opt}(\varphi)$ fraction of φ ’s clauses for every 3CNF instance φ . More specifically, the algorithm A outputs some value such that:

$$(1.1) \quad \beta \cdot \text{Opt}(\varphi) \leq A(\varphi) \leq \text{Opt}(\varphi)$$

If $\text{Opt}(\varphi) = 1$, then φ is said to be *satisfiable*.

A canonical example of a simple approximation algorithm for Max3CNF is the following scheme: for every variable in φ choose with uniform probability an assignment from $\{0, 1\}$. The probability of any given clause being satisfied by such a random assignment is $\frac{7}{8}$. Thus

$$\mathbb{E}_{x \in \{0, 1\}^n} [\# \text{ satisfied clauses of } \varphi(x)] = \frac{7m}{8}$$

where m is the number of satisfiable clauses of φ , showing that there must exist an assignment which satisfies at least $\frac{7}{8}$ of φ ’s clauses. We can derandomize this algorithm to output such an assignment. This gives us a $\frac{7}{8}$ -approximation algorithm for Max3CNF. Here’s another example:

Example 1.2. Let MaxE3Lin define the problem of the following form: let ξ be defined as a system of linear equations over the field \mathbb{F}_2 where every equation contains *exactly* 3 variables from a set of n variables x_1, \dots, x_n . Find an assignment to the variables which maximizes the number of satisfied linear equations in ξ . An example is the following system over variables x_1, x_2, x_3, x_4 :

$$\begin{aligned} x_1 + x_2 + x_3 &= 0 \\ x_1 + x_2 + x_4 &= 1 \\ x_1 + x_3 + x_4 &= 1 \end{aligned}$$

The E3Lin expression represents a linear system of Exactly 3 variables. We can extend the notation treated above for Max3CNF to this situation. In other words, if $\rho_{\xi, x}$ is the fraction of satisfied linear constraints of

ξ under assignment $x \in \mathbb{F}_2^n$.

$$\text{Opt}(\xi) = \max_{x \in \{0,1\}^n} \rho_{\xi,x}$$

A comment regarding this problem: if the given E3Lin instance has a guaranteed solution i.e $\text{Opt}(\varphi) = 1$, then Gaussian Elimination will always output an assignment which satisfies all linear constraints in polynomial-time. It is more interesting to consider instances where no such solution satisfying all of the constraints exists over \mathbb{F}_2^n . In these cases, $\text{Opt}(\xi) = 1 - \epsilon$ for some $0 < \epsilon \leq 1$. This problem also has a fairly simple $\frac{1}{2}$ -approximation algorithm: set all $x_1 = \dots = x_n = 0$ or $x_1 = \dots = x_n = 1$ depending on which satisfies the most constraints. This scheme must satisfy at least $\frac{1}{2}$ of the constraints for any instance ξ .

Example 1.3. The problem MaxCut will be casted as follows: given an undirected graph $G = (V, E)$, find the largest cut of G . There is a straightforward LP formulation for this problem:

$$(1.2) \quad \max \sum_{(u,v) \in E} e_{u,v}$$

$$(1.3) \quad e_{u,v} \leq \min \begin{cases} x_u + x_v \\ 2 - (x_u + x_v) \end{cases}$$

$$(1.4) \quad e_{u,v} \in [0, 1]$$

$$(1.5) \quad x_v \in [0, 1]$$

where the $e_{u,v}$ variables represent the edges $e \in E$ of the input graph and the x_v represent the vertex variables $v \in V$. This will actually be a relaxation from its respective integer program in which the variables can take the integer values:

$$(1.6) \quad e_{u,v} = \begin{cases} 1 & \text{if } e_{u,v} \in \mathcal{C} \\ 0 & \text{otherwise} \end{cases} \quad x_v = \begin{cases} 0 & \text{if } v \text{ is in partition } S \\ 1 & \text{otherwise} \end{cases}$$

where $\mathcal{C} \subseteq E$ is the subset of edges constituting the cut and (S, \bar{S}) is the partition defining the cut. Naturally, a solution to the integer program will also be a solution to the LP above, which implies that

$$\text{Opt}_{LP} \geq \text{MaxCut}(G)$$

There actually exists an Semi-definite Programming (SDP) relaxation for MaxCut called the *Goemans-Williamson algorithm* which will be treated in Section 6.2.

1.2. Constraint Satisfaction Problems. Generally speaking, it's useful to concretize these problems into an unified framework. The examples presented in the last section have some common interpretation shared between them: they could all be seen as manifestations of Constraint Satisfaction Problems (CSP):

Definition 1.4. Let Ω be a finite set deemed as the *domain*. A constraint satisfaction problem (CSP) Ψ over domain Ω is a finite set of predicates $\psi : \Omega^r \rightarrow \{0, 1\}$ where r is the *arity* of predicate ψ . The predicates can be of different arities. The *arity* of the CSP Ψ would be the maximum of all the arities of the predicates in Ψ .

An *instance* \mathcal{I} over CSP Ψ is a set of tuples (S, ψ) where if r is the arity of predicate ψ , $S = (v_1, \dots, v_r)$ is some ordered tuple of variables taken from finite set V consisting of variables contained in CSP Ψ . These tuples are called the *constraints* of \mathcal{I} . In addition, we add the condition that every variable show up in at

least one constraint. Variables which do not appear in any of the constraints can simply be removed.

Now given some instance \mathcal{I} , an *assignment* $F : V \rightarrow \Omega$ is simply some map between the variables and the domain. We say that F satisfies a constraint (S, ψ) if $\psi(F(S)) = 1$. For tuple $S = (v_1, \dots, v_r)$, we define $F(S) = (F(v_1), \dots, F(v_r))$. Consider the fraction of constraints of satisfied by F in \mathcal{I} :

$$(1.7) \quad \text{Val}_{\mathcal{I}}(F) = \mathbb{E}_{(S, \psi) \sim \mathcal{I}}[\psi(F(S))]$$

By taking the maximum fraction over all such assignments:

$$(1.8) \quad \text{Opt}(\mathcal{I}) = \max_{F: V \rightarrow \Omega} \text{Val}_{\mathcal{I}}(F)$$

Example 1.5.

- For **Max3CNF**, the domain would be $\Omega = \{0, 1\}$ and the CSP Ψ would be composed of the single predicate $\vee_3 : \{0, 1\}^3 \rightarrow \{0, 1\}$. This predicate is just the logical OR of the three input variables. Any 3CNF φ can have its clauses be expressed as constraints (S, \vee_3) where S would be the variables inputted into \vee_3 . Hence, an assignment F would be an assignment into the variables found in 3CNF φ , showing that $\text{Opt}(\varphi)$ aligns with the definition given in the last section.
- For **MaxE3Lin**, the domain would be $\Omega = \mathbb{F}_2$ and Ψ consist of predicates of the form $(x_1, \dots, x_3) = x_1 + x_2 + x_3$ and $(x_1, \dots, x_3) = x_1 + x_2 + x_3 + 1$ representing both types of linear constraints found in a system of three-variable equations over \mathbb{F}_2 . An instance ξ would consist of constraints (S, ψ) where S would be a three variable tuple containing the variables appearing in the linear constraint ψ .
- For **MaxCut**, the domain can be defined as $\Omega = \{-1, 1\}$ with Ψ set to the simple predicate $\neq : \{-1, 1\}^2 \rightarrow \{0, 1\}$. This simply tests if the inputted values are not equal. The variable set V of an instance \mathcal{I} would be indexed by the vertices contained in a graph $G = (V, E)$. There is one constraint tuple, $((v_i, v_j), \neq)$ for every edge $(v_i, v_j) \in E$. Thus, an assignment $F : V \rightarrow \{-1, 1\}$ would assign the vertices into two partitions such that a constraint is satisfied iff the corresponding edge is contained in \mathcal{C} .
- The CSP for **Max3Color**, the maximization counterpart for the NP-complete decision problem, **3Color**, is similarly defined to that of **MaxCut** except our domain would be $\Omega = \{0, 1, 2\}$. This signifies the three possible colors to color any vertex $v \in V$ in an input graph $G = (V, E)$.

As implied in the examples above, there is a generic method to formulate a maximization problem in respect to a given CSP Ψ :

Definition 1.6. For a given CSP Ψ , formulate $\text{MaxCSP}(\Psi)$ as the problem: given an instance \mathcal{I} , output an assignment F which satisfies the largest number of constraints in \mathcal{I} .

Note that Definition 1.1 for a β -approximation algorithm extends to this maximization problem in the natural way.

1.3. Gap Problems. The NP-hardness theory frequently relies on Karp reductions from decision problems to decision problems. In light of this, we can tailor this paradigm to optimization problems in the form of so-called *gap problems*.

Definition 1.7. A *promise problem* is defined as a tuple (YES, NO) where $\text{YES}, \text{NO} \subseteq \Sigma^*$ in respect to some alphabet Σ . Furthermore, we require that $\text{YES} \cap \text{NO} = \emptyset$ but not necessarily that $\text{YES} \cup \text{NO} = \Sigma^*$.

Definition 1.8. Given a $\text{MaxCSP}(\Psi)$ problem, we define $\text{Gap}_{\alpha,\beta}\text{MaxCSP}(\Psi)$ for $\alpha < \beta$ as the promise problem: given an instance \mathcal{I} :

$$(1.9) \quad \mathcal{I} \in \text{YES} \iff \text{Opt}(\mathcal{I}) \geq \beta$$

$$(1.10) \quad \mathcal{I} \in \text{NO} \iff \text{Opt}(\mathcal{I}) < \alpha$$

Furthermore, an algorithm A decides $\text{MaxCSP}(\Psi)$ if for input instance \mathcal{I} it accepts iff $\mathcal{I} \in \text{Yes}$ and rejects iff $\mathcal{I} \in \text{No}$. If $\mathcal{I} \notin \text{Yes} \cup \text{No}$, we do not care what the algorithm outputs.

In particular, we deem a $\text{Gap}_{\alpha,\beta}\text{MaxCSP}(\Psi)$ problem as NP-hard if for every language $L \in \text{NP}$, there exists a polynomial-time reduction f taking input strings $x \in \{0,1\}^*$ to CSP Ψ instances such that:

$$x \in L \implies \text{Opt}(f(x)) \geq \beta$$

$$x \notin L \implies \text{Opt}(f(x)) < \alpha$$

The NP-hardness of approximation algorithms reduces to that of gap problems as shown in the below observation:

Theorem 1.9. Suppose $\text{Gap}_{\alpha,\beta}\text{MaxCSP}(\Psi)$ is NP-hard for CSP Ψ , then approximating $\text{MaxCSP}(\Psi)$ to at least an $\frac{\alpha}{\beta}$ factor is NP-hard.

Proof. Suppose there exists an algorithm A which can $\frac{\alpha}{\beta}$ -approximate $\text{MaxCSP}(\Psi)$. For an instance \mathcal{I} such that $\text{Opt}(\mathcal{I}) \geq \beta$:

$$A(\mathcal{I}) \geq \frac{\alpha}{\beta} \cdot \text{Opt}(\mathcal{I}) = \frac{\alpha}{\beta} \cdot \beta = \alpha$$

Else if $\text{Opt}(\mathcal{I}) < \alpha$

$$A(\mathcal{I}) \leq \text{Opt}(\mathcal{I}) < \alpha$$

by Definition 1.1 adapted to $\text{MaxCSP}(\Psi)$ instances. Hence, the algorithm can decide $\text{Gap}_{\alpha,\beta}\text{MaxCSP}(\Psi)$ by checking its outputted value in respect to α . \square

This in particular demonstrates that showing the hardness of approximating a particular problem is equivalent to showing the hardness of its corresponding gap problem.

2. THE PCP THEOREM

2.1. Intuitions. This section will be centered around the seminal PCP Theorem [2], [3], which characterized NP in a framework considered unconventional at the time. PCPs, or Probabilistically Checkable Proofs, represent a twist on the idea of NP. Recall that NP roughly represents the languages which have verifiers which can check proofs of membership in polynomial time. PCPs represent an extension of this definition where the verifier can be *probabilistic* and is granted *random access* to the proof string π . If we allow the verifier to simply query π by outputting a index i , it has access to $\pi[i]$. Since we can express an index in $\log n$ bits, this in theory gives the verifier access to proof strings of exponential length. To formalize these notions, we begin with definitions:

Definition 2.1. Given a language L and $r, q : \mathbb{N} \rightarrow \mathbb{N}$, a $(r(n), q(n))$ -PCP-verifier for L consists of a polynomial-time algorithm V with the following properties:

- For input strings $x \in \{0, 1\}^n$, $\pi \in \{0, 1\}^{\leq N}$ for $N = q(n)2^{r(n)}$, V makes $r(n)$ coin flips and decides $q(n)$ queries addresses $i_1, \dots, i_{q(n)}$ of the proof π . Based on these queries, it outputs 1 for “accept” or 0 for “reject”.
- (Completeness) For $x \in L$, there exists some proof π such that $V(x, \pi, r) = 1$ for all random coin tosses r . In other words:

$$(2.1) \quad \mathbb{P}_r[V(x, \pi, r) = 1] = 1$$

- (Soundness) For $x \notin L$, for all proofs π :

$$(2.2) \quad \mathbb{P}_r[V(x, \pi, r) = 1] \leq \frac{1}{2}$$

Define the class $\text{PCP}(r(n), q(n))$ as the set of languages L which has a $(c \cdot r(n), d \cdot q(n))$ -PCP-verifier for some constants $c, d > 0$.

Remark. Sometimes the completeness criterion is too strong for our purposes (see the comments on **MaxE3Lin** problem). In these cases, we like to denote the class $\text{PCP}_{\alpha, \beta}(r(n), q(n))$ as the languages L which have a $(r(n), q(n))$ -NP-verifier such that the completeness and soundness criteria are amended as below:

- (Completeness) For $x \in L$, there exists some proof π such that $V(x, \pi, r) = 1$ for all random coin tosses r . In other words:

$$(2.3) \quad \mathbb{P}_r[V(x, \pi, r) = 1] \geq \beta$$

- (Soundness) For $x \notin L$, for all proofs π :

$$(2.4) \quad \mathbb{P}_r[V(x, \pi, r) = 1] < \alpha$$

Here, β is the *completeness parameter* while α is the *soundness parameter*. The class introduced in the original definition would thus be denoted as $\text{PCP}_{\frac{1}{2}, 1}(r(n), q(n))$. PCP verifiers whose completeness parameter is one ($\beta = 1$) is deemed as *perfectly complete*.

The PCP Theorem says that **NP** is *exactly* the class of PCPs which uses a *logarithmic* number of random bits and a *constant* number of queries.

Theorem 2.2. (*The PCP Theorem* [2], [3])

$$(2.5) \quad \text{NP} = \text{PCP}_{\frac{1}{2}, 1}(O(\log n), O(1))$$

Actually, one direction of this theorem is not too difficult to see:

Proposition 2.3. For every constants $Q \in \mathbb{N}, c > 0$, $\text{PCP}_{\frac{1}{2}, 1}(c \cdot \log n, Q) \subseteq \text{NP}$

Proof. Begin with the observation that $\text{PCP}_{\frac{1}{2}, 1}(r(n), q(n)) \subseteq \text{NTIME}(q(n)2^{r(n)})$. This is justified by the view of an **NTIME** machine simulating the verifier by trying all possible coin tosses and queries to the input string x and proof string π . It can then count all of the accepting paths to determine the probability of acceptance. If $q = O(1)$ and $r = O(\log n)$, then the right side of the inclusion will be $\text{NTIME}(2^{O(\log n)}) = \text{NP}$. Here we use the definition: $\text{NP} = \bigcup_{c \in \mathbb{N}} \text{NTIME}(n^c)$ \square

Remark. The set of queries a PCP verifier makes could be either *adaptive* or *non-adaptive*. Adaptive queries can be dependent on the outcome of previous queries while non-adaptive queries are independent of one another. The verifiers in these notes will all be non-adaptive for the sake of presentation. The PCP Theorem still holds when the verifier makes adaptive queries. The only change would be that the proof length would be at most $2^{r(n)+q(n)}$ rather than at most $q(n)2^{r(n)}$.

2.2. Equivalence of PCP Theorems. It may be difficult to understand the importance of the PCP Theorem in its form presented in Theorem 2.2. It turns out there are other equivalent forms of the PCP Theorem more palatable in the context of our goal to prove hardness-of-approximation results.

Theorem 2.4. (PCP Theorem: Gap3SAT-hardness) *The problem $\text{Gap}_{\alpha,1}\text{Max3SAT}$ is NP-hard. In other words there exists a constant $\alpha < 1$ such that, for every NP language L , there exists a polynomial-time reduction f mapping L to 3SAT formulas such that:*

$$\begin{aligned} x \in L &\implies \text{Opt}(f(x)) = 1 \\ x \notin L &\implies \text{Opt}(f(x)) \leq \alpha \end{aligned}$$

An immediate consequence of Theorem 2.4 and Theorem 1.9 is that if there exists an α -approximation algorithm for Max3SAT, then $P = NP$. With this, we have the first steps towards an inapproximability result: if $P \neq NP$, there exists no efficient Max3SAT algorithm which can approximate better than an α factor. Note that we haven't actually found a concrete value for α yet. This will be addressed once we prove Håstad's 3-bit PCP for NP in a future section.

Theorem 2.5. (PCP-Theorem: GapMaxqCSP-hardness) *There exists constants $q \in \mathbb{N}$ and $\alpha < 1$ where the problem $\text{Gap}_{\alpha,1}\text{MaxqCSP}$ is NP-hard. To elaborate, for every NP language L , there exists a polynomial time reduction mapping an L to a instance $f(x)$ of some CSP Ψ over domain $\Omega = \{0,1\}$ where Ψ consists of q -ary predicates, such that*

$$\begin{aligned} x \in L &\implies \text{Opt}(f(x)) = 1 \\ x \notin L &\implies \text{Opt}(f(x)) \leq \alpha \end{aligned}$$

Theorem 2.6. *All the PCP Theorems above are equivalent to one another.*

Before we embark on the proof, let us establish an equivalence between PCPs and CSPs:

Lemma 2.7. (Equivalence between PCPs and CSPs) *Theorem 2.2 and Theorem 2.5 are equivalent.*

Proof. First, assume $NP = \text{PCP}_{\frac{1}{2},1}(O(\log n), O(1))$. We will outline a procedure to convert the verifier V into an instance \mathcal{I} for a q -ary CSP Ψ for some constant q . For some input string $x \in \{0,1\}^n$ and proof string π , let $r \in \{0,1\}^{c \cdot \log n}$ be the coin flips made by V and $V_{x,r}$ be the deterministic procedure which is executed on input x and coin flip r such that $V_{x,r} = 1$ iff V accepts proof π on input x and coin flip r . We can define the domain of our constructed CSP Ψ to be $\Omega = \{0,1\}$ and the predicates to be $\{V_{x,r}\}_r$. Now our instance \mathcal{I} of Ψ is casted as the tuples $(S, V_{x,r})$ where S will be at most a q -sized tuple indicating which indices of the proof π are queried when conditioned on r . This yields a polynomially-sized qCSP instance \mathcal{I} . Furthermore, since the verifier V runs in polynomial time, its execution can be simulated on all r to output the instance \mathcal{I} in polynomial-time. Thus, we have given a polynomial-time reduction from an input x to its corresponding CSP instance \mathcal{I} , so Theorem 2.5.

Conversely, suppose we had a reduction from NP to $\text{Gap}_{\alpha,1}\text{MaxqCSP}$ as stated in Theorem 2.5. We devise polynomial-time reduction taking an instance MaxqCSP to a polynomial-time PCP verifier V using logarithmic number of random bits and a constant number of queries to the supplied proof π . For an input $x \in \{0,1\}^n$, the proof will be expected to be an assignment to its respective instance $f(x)$ in the notation utilized in Theorem 2.5. Verifier V makes coin flips $r \in \{0,1\}^{c \cdot \log n}$ to choose one constraint tuple (S, ψ) where ψ is some q -ary predicate. Only a logarithmic number of random bits are required to query any constraint in instance $f(x)$ as the polynomial-time NP reduction can only generate a polynomial number

of such constraints. The PCP only has to make q -queries to the proof π to find the assignments to the variables listed in $S = (v_1, \dots, v_q)$. By the properties listed in Theorem 2.5, the PCP verifier V must have completeness 1 and soundness $\leq \frac{1}{2}$ as claimed. \square

Lemma 2.8. *(Equivalence between GapMaxqCSP and GapMax3SAT Theorem 2.4 and Theorem 2.5 are equivalent.)*

Proof. Since any 3SAT instance can be seen as a particular type of 3CSP instance, one direction is immediate. Conversely, if we assume Theorem 2.5, we aim to find some constant α' such that there exist a reduction from an instance of Gap $_{1-\alpha,1}$ MaxqCSP for $q \in \mathbb{N}$ claimed in Theorem 2.5 to an instance of Gap $_{1-\alpha',1}$ Max3SAT. Let the CSP Ψ be comprised q -ary predicates ψ . For an instance \mathcal{I} of Ψ , a constraint tuple (S, ψ) can be expressed as a logical AND of 2^q clauses where each clause is a logical OR of q variables or their negations. In other words, (S, ψ) is essentially a CNF of width q and of size at most 2^q . We can then construct an “equivalent” 3CNF ψ'_S as follows: add extra symbols $\Pi_1, \dots, \Pi_{(q-3)2^q}$. It can be shown that there exists a 3CNF ψ'_S of size at most $(q-2)2^q$ such that:

- (1) For every $x \in \{0, 1\}^q$ which causes $\psi(x) = 1$, there exists an assignment Π such that $\psi'(x, \Pi) = 1$.
- (2) Else if $\psi(x) = 0$, then for all assignments Π , $\psi'(x, \Pi) = 0$

Now take the total 3CNF defined by a conjunction of all such ψ'_S :

$$\psi_{\mathcal{I}} = \bigwedge_{(S, \psi) \in \mathcal{I}} \psi'_S$$

The total formula $\psi_{\mathcal{I}}$ is determined by at most $m q 2^q$ number of clauses and at most $n + m(q-3)2^q$ number of variables. If an instance \mathcal{I} has all of its constraints satisfiable by some assignment, then by property (1) listed above, there must exist an assignment to the variables of the $\psi_{\mathcal{I}}$ such that it is satisfied. On the other hand, if for all assignments to instance \mathcal{I} only satisfy at most an $1 - \alpha$ fraction of constraints, then fraction of clauses satisfied can be at most $1 - \alpha + \alpha(1 - \frac{1}{(q-2)2^q}) = 1 - \frac{\alpha}{(q-2)2^q}$. This is due to the fact that for each unsatisfied ψ'_S , the most number of its clauses which can be satisfied by any assignment will be $(q-2)2^q - 1$. Taking $\alpha' = \frac{\alpha}{(q-2)2^q}$ yields the required constant. \square

Remark. The existence of the “equivalent” 3CNF is outlined in Problem 7.11 of the O’Donnell textbook [15].

Remark. The proof of the PCP Theorem (Theorem 2.2) is quite involved and out of the scope of these notes. The original proof hinged on clever algebraic techniques such as low-degree testing [3]. A simpler combinatorial proof by Dinur arrived years later [5]. For more information, refer to the cited publications or Chapter 22 of the Arora, Barak textbook [1].

3. LABEL-COVER AND PROJECTION GAMES

We now introduce a problem which manages to provide a natural paradigm for capturing the essence of CSPs and proving inapproximability results. These “projection games” were introduced by Bellare, Goldreich, and Sudan [4]. The NP-hardness of the gap problem version of Label Cover was used by Håstad to show tight inapproximability results for Max3SAT and MaxE3Lin [10].

Definition 3.1. A *Label Cover (LC) Problem* instance \mathcal{G} is defined by a bipartite graph $(A \sqcup B, E)$, finite alphabets Σ_A, Σ_B , and a set of projections $\pi_e : \Sigma_A \rightarrow \Sigma_B$ for every edge $e \in E$. Define an *assignment* as consisting of two maps $\mathfrak{A} : A \rightarrow \Sigma_A$, $\mathfrak{B} : B \rightarrow \Sigma_B$. An edge $e = (a, b) \in E$ is said to be satisfied by this assignment if the assignment is compatible with projection π_e :

$$(3.1) \quad \pi_e(\mathfrak{A}(a)) = \mathfrak{B}(b)$$

The value of this game will be

$$(3.2) \quad \text{Opt}(\mathcal{G}) = \max_{(\mathfrak{A}, \mathfrak{B})} \mathbb{E}_{e \sim E} [e \text{ satisfied}]$$

In other words, the value will be the largest fraction of edges satisfied by any assignment to the vertices. The corresponding gap problem for Label Cover, $\text{Gap}_{\alpha, \beta} \text{LC}$, is defined as the promise problem:

$$\begin{aligned} \text{YES} &= \{\mathcal{G} \mid \text{Opt}(\mathcal{G}) \geq \beta\} \\ \text{NO} &= \{\mathcal{G} \mid \text{Opt}(\mathcal{G}) < \alpha\} \end{aligned}$$

In the case of perfect completeness, we abbreviate $\text{Gap}_{\alpha, 1} \text{LC}$ as simply $\text{Gap}_{\alpha} \text{LC}$.

There are a few observations worthy of mentioning here. The first regards a type of equivalence between CSP instances and Label Cover instances. Specifically, let \mathcal{I} be an instance of a given r -ary CSP Ψ over domain Ω . We can translate this CSP instance into a Label Cover instance as follows: Let the left-hand partition A of our bipartite graph be indexed by the set of constraints (S, ψ) and the right-hand partition B be indexed by the variables of the CSP V . Draw an edge from a constraint tuple (S, ψ) to a variable v if that variable appears in S . Set $\Sigma_A = \{(q_1, \dots, q_r) \in \Omega^r \mid \exists \psi \in \Psi, \psi(q_1, \dots, q_r) = 1\}$ and $\Sigma_B = \Omega$, and for every edge $e = ((v_1, \dots, v_r), \psi), v)$ define the projection $\pi_e : \Sigma_A \rightarrow \Sigma_B$ to be

$$\pi_e(\omega_1, \dots, \omega_r) = \omega_i \text{ if } v_i = v$$

On the other hand, every Label Cover instance can be seen as a 2CSP over a sufficiently large domain: the predicates of the CSP would be all 2-ary predicates $\pi : \Sigma_A \times \Sigma_B \rightarrow \{0, 1\}$ representing every possible map from $\Sigma_A \rightarrow \Sigma_B$. Thus, the domain of our CSP can be defined as $\Omega = \Sigma_A \cup \Sigma_B$. The corresponding instance of this CSP would be (S, π_e) where π_e represents the predicate corresponding to the edge e 's projection map π_e and $S = (a, b)$ would be the vertices of e between A and B respectively.

By Theorem 2.7 and the observations made above, there is a method of converting the PCP for NP given by the PCP Theorem to a CSP, then converting that CSP into a Label Cover instance. This shows that $\text{Gap}_{\alpha} \text{LC}$ for some $\alpha > 0$ must be NP-hard:

Theorem 3.2. (*Weak Projection Games Theorem*) $\text{Gap}_{\alpha} \text{LC}$ for some $\alpha > 0$ is NP-hard.

The primary drawback of this theorem is that it does give us an NP-hardness result for *every* constant $\alpha > 0$. This will be addressed by Raz's Parallel Repetition Theorem.

3.1. Raz's Parallel Repetition Theorem. In this section, we give a cursory outline of Raz's Parallel Repetition Theorem and its consequences. First, we can extend Definition 2.1 to include proof strings over non-boolean alphabets. To accomodate such verifiers, we can extend the definition to $\text{PCP}_{\beta, \alpha}^{\Sigma}(r(n), q(n))$, the class of languages with polynomial-time verifiers taking proof strings π over non-boolean alphabet Σ and uses $r(n)$ random bits and at most $q(n)$ queries to π . The following theorem shows that one can reduce the number of queries to two at the expense of a weaker soundness constant and a larger alphabet:

Theorem 3.3. $\text{PCP}_{1-\alpha, 1}^{\Sigma}(r(n), q(n)) \subseteq \text{PCP}_{1-\frac{\alpha}{q}, 1}^{\Sigma^q}(r(n) + \log q(n), 2)$

Proof. Given a verifier for a language $L \in \text{PCP}_{\beta, 1-\alpha}^{\Sigma}(r(n), q(n))$, it behaves as such on input $x \in \{0, 1\}^n$:

- (1) Flips a $r(n)$ coins which we denote as R .
- (2) Using R , it computes a series of indices i_1, \dots, i_q where $q = q(n)$.

- (3) It queries $\pi_{i_1}, \dots, \pi_{i_q} \in \Sigma$ from the proof π .
- (4) Finally, it feeds the symbols in a predicate $V_{x,R}(\pi_{i_1}, \dots, \pi_{i_q})$ which outputs “accept” or “reject”.

We extend this to a modified verifier V' to accomodate the concatenation of two proof strings: if $m = |\Sigma|$ $\pi_1 : [m]^q$ then $\rightarrow \Sigma^q$ and $\pi_2 : [m] \rightarrow \Sigma$:

- (1) Flips a $r(n)$ coins which we denote as R .
- (2) Using R , it computes a series of indices i_1, \dots, i_q where $q = q(n)$
- (3) Queries $\pi' = \pi_1(i_1, \dots, i_q)$
- (4) Computes the predicate:

$$V'_{x,R}(\pi') = V_{x,R}(\pi'_1, \dots, \pi'_q)$$

- (5) Finally, chooses random $\ell \in [q]$ and checks if $\pi_2(i_\ell) = \pi'_\ell$.
- (6) The verifier accepts iff both checks pass.

For completeness, if there exists a proof π which the original verifier V accepts with probability one, then setting $\pi_1(i_1, \dots, i_q) = (\pi_{i_1}, \dots, \pi_{i_q})$ and $\pi_2 = \pi$ will induce V' to accept with probability one by construction. As for soundness, suppose at least α fraction of the predicates $V_{x,R}$ reject in respect to some proof string π . Each one of the predicates contained in this fraction must fail the check at Step 4 if $\pi' = \pi_1(i_1, \dots, i_q) = (\pi_{i_1}, \dots, \pi_{i_q})$. However, π_1 does not have to respect this rule for any q -tuple of queries (i_1, \dots, i_q) given as input. If so, there must exist at least one index ℓ such that $\pi'_\ell \neq \pi_2(i_\ell) = \pi_\ell$. Since such an index is chosen with probability $\frac{1}{q}$, we have that V' must reject with probability at least $\frac{\alpha}{q}$. This completes the proof. \square

Raz’s Parallel Repetition Theorem constructed a verifier which allowed the soundness constant to drop exponentially with the cost of bloating the alphabet size:

Theorem 3.4. (*Raz’s Parallel Repetition Theorem* [16]) *For all $s \in (0, 1)$, there exists $c_s \in (0, s)$ such that:*

$$\text{PCP}_{s,1}^\Sigma(r, 2) \subseteq \text{PCP}_{c_s,1}^{\Sigma^t}(rt, 2)$$

Corollary 3.4.1. *For all $\epsilon > 0$, there exists an alphabet Σ such that $|\Sigma| \leq \text{poly}(\frac{1}{\epsilon})$:*

$$\text{NP} \subseteq \text{PCP}_{\epsilon,1}^\Sigma(O(\log n \cdot \log 1/\epsilon), 2)$$

Proof. For some constant $Q > 0$,

$$\begin{aligned} \text{NP} &\subseteq \text{PCP}_{\frac{1}{2},1}^{\{0,1\}}(O(\log n), Q) \\ &\subseteq \text{PCP}_{1-\frac{1}{2Q},1}^{\{0,1\}^Q}(O(\log n), 2) \\ &\subseteq \text{PCP}_{c_Q^t,1}^{\{0,1\}^{Q^t}}(O(t \log n), 2) \\ &\subseteq \text{PCP}_{\epsilon,1}^\Sigma(O(\log n \cdot \log 1/\epsilon), 2) \end{aligned}$$

where $|\Sigma| \leq \text{poly}(\frac{1}{\epsilon})$. The second inclusion follows from Theorem 3.3 and the third inclusion follows from Theorem 3.4. \square

By combining Corollary 3.4.1, the NP-hardness of Label Cover comes to fruition:

Theorem 3.5. (*Projection Games Theorem*) *For every $\epsilon > 0$, there exist alphabets Σ_A, Σ_B where $|\Sigma_A|, |\Sigma_B| \leq \text{poly}(\frac{1}{\epsilon})$ such that $\text{Gap}_\epsilon \text{LC}$ is NP-hard.*

This is shown by embedding the query pairs of the verifier into the projection constraints of a Label Cover instance graph.

4. HÅSTAD'S 3-BIT PCP

4.1. Inapproximability Results for MaxE3Lin and Max3Sat. The significance of Håstad's PCP for NP is its role in showing tight inapproximability results for MaxE3Lin and hence Max3Sat. The theorem is first stated below:

Theorem 4.1. (*Håstad's 3-bit PCP*, [10]) *For every $\delta > 0$ and $L \in \text{NP}$, there exists a PCP verifier for L over the boolean alphabet such that for every*

- *The verifier V queries 3 bits of the proof $x_{q_1}, x_{q_2}, x_{q_3} \in \pi$ such that verification predicate is a three variable linear equation over \mathbb{F}_2 depending on the queried bits $x_{q_1}, x_{q_2}, x_{q_3}$.*
- *If $x \in L$, then there exists a proof π :*

$$(4.1) \quad \mathbb{P}[V(x, \pi, r) = 1] \geq 1 - \delta$$

- *If $x \notin L$, then for all proofs π :*

$$(4.2) \quad \mathbb{P}[V(x, \pi, r) = 1] \leq \frac{1}{2} + \delta$$

On closer inspection, Theorem 4.1 reveals that a gap problem based on MaxE3Lin is actually NP-hard to decide:

Corollary 4.1.1. *For every $\delta > 0$, the promise problem $\text{Gap}_{\frac{1}{2}+\delta, 1-\delta}\text{MaxE3Lin}$ is NP-hard.*

This gives our first threshold result: Corollary 4.1.1 states that the simple $\frac{1}{2}$ -approximation algorithm introduced in Example 1.2 is optimal in the sense that the existence of an efficient algorithm improving on the $\frac{1}{2}$ constant factor would show that $\text{P} = \text{NP}$. From this result, by reducing MaxE3Lin to Max3Sat, another threshold result appears:

Corollary 4.1.2. *For every $\delta > 0$, the promise problem $\text{Gap}_{\frac{7}{8}+\delta, 1-\delta}\text{Max3Sat}$ is NP-hard.*

Proof. It suffices to transform MaxE3Lin CSP instances to MaxSat CSP instances. For predicates of the form $x + y + z = 1$, consider the conjunction of the clauses:

$$\begin{aligned} x \vee y \vee z \\ \bar{x} \vee y \vee \bar{z} \\ \bar{x} \vee \bar{y} \vee z \\ x \vee \bar{y} \vee \bar{z} \end{aligned}$$

The case for predicates of the form $x + y + z = 0$ are handled with a similar idea. For an assignment of a MaxE3Lin instance which satisfies a linear constraint, it must satisfy all four of the associated MaxSat constraints. In the case where an assignment does not satisfy a linear constraint, at most three of the four clauses will be satisfied. Hence if an assignment satisfies $\frac{1}{2} + \delta$ of the constraints, then it can satisfy at most

$$\left(\frac{1}{2} + \delta\right) 4 + \left(\frac{1}{2} - \delta\right) 3 = \frac{7}{2} + \frac{3\delta}{2}$$

fraction of clauses. On the other hand, if at least $1 - \delta$ of the linear constraints are satisfied then at least $1 - \delta$ of the clauses must also be satisfied. \square

The proof of Theorem 4.1 leverages the NP-hardness of $\text{Gap}_\epsilon \text{LC}$ mentioned in Theorem 3.5. Håstad devised a dictatorship test based on E3Lin predicates then composed this test with checking the validity of a $\text{Gap}_\epsilon \text{LC}$ proof. In fact, the three queries to the proof arise from a modified Blum-Luby-Rubinfeld Linearity Test. The details will be treated in the upcoming subsections.

4.2. The Long Code. The workhorse for Håstad's proof will be the use of the *long code*. A long code is essentially the truth table of a dictator function. The idea roughly revolve around the idea that dictator functions are locally-testable as linear functions (à la BLR Test) and have an useful error-correction property. Correct proofs will be expected to encode labels to vertices in a Label Cover instance as long codes of dictators corresponding to the indices of the labels. In other words, for a label $\ell \in \Sigma$, a corresponding index $i(\ell)$ can be computed with $\log |\Sigma|$ bits. The total number of such bitstrings will be, of course, length $2^{\log |\Sigma|} = |\Sigma|$. We define the long code of ℓ as the bitstring:

$$(4.3) \quad \text{long}_\ell = (f(\ell) \mid f : \Sigma \rightarrow \{-1, 1\})$$

If we additionally encode this bitstring into the dictator function $\chi_{i(\ell)}(x_1, \dots, x_{|\Sigma|}) = x_{i(\ell)}$, the truth table for this function will be of total length $2^{|\Sigma|}$. This truth table will exactly be long_ℓ . Here, the input into $\chi_{i(\ell)}$ will be the truth tables for the individual boolean functions $f : \Sigma \rightarrow \{-1, 1\}$. This reveals that the long code requires a *double-exponential* proof length. The truth table for $\chi_{i(\ell)}$ will serve as query access to the modified BLR test mentioned above.

4.3. Aside on Dictatorship Testing. The first step towards testing if an input boolean function $f : \{-1, 1\}^n \rightarrow \{-1, 1\}$ is a dictator arises from the Blum-Luby-Rubinfeld Test (BLR) which we restate below:

- (1) Sample $x, y \sim_R \{-1, 1\}^n$
- (2) Accept iff $f(x)f(y) = f(xy)$

Theorem 4.2. *Suppose the BLR test accepts $f : \{-1, 1\}^n \rightarrow \{-1, 1\}$ with probability $1 - \epsilon$, then f is ϵ -close to a linear function χ_{S^*} for some $S^* \subseteq [n]$.*

Certainly the dictators χ_i , $i \in [n]$ are linear functions. Hence, if $f = \chi_i$ for some i , then the BLR test accepts with probability one. However, we have the rest of the parity functions χ_S , $|S| \geq 2$ which the BLR Test cannot distinguish. We need to amend the test to ensure that parity functions of higher weight are rejected with high probability. Håstad proposed modifying the vanilla BLR test to add noise to the sampled product xy . Although this sacrifices perfect completeness, it penalizes large parity functions:

- (1) Sample $x, y \sim \{-1, 1\}^n$
- (2) Sample $z \sim N_{1-2\epsilon}(xy)$
- (3) Accept iff $f(x)f(y) = f(z)$

Lemma 4.3. *(Completeness of the Noisy BLR Test) If $f = \chi_i$ is a dictator for some i , then*

$$\mathbb{P}[\text{Noisy BLR test accepts}] \geq 1 - \epsilon$$

Proof. Note that if $z \sim N_{1-2\epsilon}(xy)$ for some $x \in \{-1, 1\}^n$, then y can be expressed as below:

$$(4.4) \quad z_i = \begin{cases} x_i y_i & \text{with probability } 1 - \epsilon \\ -x_i y_i & \text{with probability } \epsilon \end{cases}$$

By the acceptance criterion,

$$f(z) = f(x)f(y) \implies z_i = x_i y_i$$

This occurs with probability $1 - \epsilon$ by the observation above, yielding completeness as claimed. \square

Lemma 4.4. *Suppose that for some constant $\nu > 0$:*

$$\mathbb{P}[\text{Noisy BLR test accepts}] \geq \frac{1}{2} + \nu$$

then

$$(4.5) \quad 2\nu \leq \sum_{S \subseteq [n]} \hat{f}(S)^3 (1 - 2\epsilon)^{|S|}$$

Proof. The proof begins by noticing the accepting probability can be expressed as:

$$\mathbb{P}[\text{Noisy BLR test accepts}] = \frac{1}{2} + \frac{1}{2} \mathbb{E}_{x,y,z} [f(x)f(y)f(z)]$$

By our assumption, we prove that:

$$\begin{aligned} \frac{1}{2} + \nu &\leq \frac{1}{2} + \frac{1}{2} \mathbb{E}_{x,y,z} [f(x)f(y)f(z)] \implies \\ 2\nu &\leq \mathbb{E}_{x,y,z} [f(x)f(y)f(z)] \implies \\ 2\nu &\leq \mathbb{E}_{x,y,z} [f(x)f(y) \mathbb{E}_{z \sim N_{1-2\epsilon}(xy)} [f(z)]] \implies \\ 2\nu &\leq \mathbb{E}_x [f(x) \mathbb{E}_y [f(y) \mathcal{T}_{1-2\epsilon} f(xy)]] \\ 2\nu &\leq \mathbb{E}_x [f(x) (f * \mathcal{T}_{1-2\epsilon} f)(x)] \implies \\ 2\nu &\leq \sum_{S \subseteq [n]} \hat{f}(S) \hat{f}(S) \widehat{\mathcal{T}_{1-2\epsilon} f}(S) \implies \\ 2\nu &\leq \sum_{S \subseteq [n]} \hat{f}(S)^3 (1 - 2\epsilon)^{|S|} \end{aligned}$$

as claimed. \square

The operator \mathcal{T}_ρ is the *noise operator* which, given boolean function $f : \{-1, 1\}^n \rightarrow \{-1, 1\}$, maps f to another boolean function $\mathcal{T}_\rho(f)$ with the following Fourier decomposition:

$$(4.6) \quad \mathcal{T}_\rho(f) = \sum_{S \subseteq [n]} \rho^{|S|} \hat{f}(S) \chi_S$$

Corollary 4.4.1. (*Soundness of the Noisy BLR Test*) *There exists some $S^* \subseteq [n]$ such that*

$$(4.7) \quad |\hat{f}(S^*)| \geq 2\nu \quad |S^*| \leq O\left(\frac{1}{\epsilon} \log \frac{1}{\nu}\right)$$

Proof. From Theorem 4.4, we deduce that:

$$\begin{aligned} 2\nu &\leq \sum_{S \subseteq [n]} \hat{f}(S)^3 (1 - 2\epsilon)^{|S|} \leq \max_S \hat{f}(S) (1 - 2\epsilon)^{|S|} \sum_{S \subseteq [n]} \hat{f}(S)^2 \\ &= \max_S \hat{f}(S) (1 - 2\epsilon)^{|S|} \end{aligned}$$

By taking the maximum attained as S^* ,

$$2\nu \leq \hat{f}(S^*) (1 - 2\epsilon)^{|S^*|}$$

Both $\hat{f}(S^*) \geq 2\nu$ and $(1 - 2\epsilon)^{|S^*|} \geq 2\nu$ (as f is boolean). By the inequality $1 - x \leq e^{-x}$,

$$e^{-2\epsilon|S^*|} \geq 2\nu$$

which shows that $|S^*| \leq O\left(\frac{1}{\epsilon} \log \frac{1}{\nu}\right)$. \square

Remark. This theorem has a similar interpretation with the Friedgut-Kalai-Naor Theorem in the rough sense that boolean functions who have their Fourier weights concentrated in the lower degrees are similar to dictator functions. Indeed this is the crux behind the soundness property of the Noisy BLR test.

4.4. The Code Consistency Test. The consistency test will harness the dictatorship test formulated above to check the validity of an alleged long codes and satisfaction of the sampled edge by the labels encoded by the long codes. Formally speaking, the input will be a $\text{Gap}_\epsilon \text{LC}$ instance $((A \sqcup B, E), \Sigma_A, \Sigma_B, \{\pi_e\}_{e \in E})$ and the proof will be the long codes encoding the assignments to the vertices, $\mathfrak{A} : A \rightarrow \Sigma_A$, $\mathfrak{B} : B \rightarrow \Sigma_B$. The consistency check will aim to perform two tasks simultaneously in respect to an edge $e = (a, b)$ and its associated projection constraint π_e : Let $f_a : \{-1, 1\}^{|\Sigma_A|} \rightarrow \{-1, 1\}$, $f_b : \{-1, 1\}^{|\Sigma_B|} \rightarrow \{-1, 1\}$ be the functions claimed to be long codes for the labels assigned to vertices a, b : $\mathfrak{A}(a), \mathfrak{B}(b)$. Recall these will be encoded as truth table strings on the proof π .

Validity Check: Check that the functions f_a, f_b are indeed valid long codes for some labels $\mathfrak{A}(a), \mathfrak{B}(b)$.

Consistency Check: Verify that the assigned labels satisfy π_e

$$\pi_e(\mathfrak{A}(a)) = \mathfrak{B}(b)$$

Before presenting the test, some notation is in order. Let $(x \circ \pi_e)_i = x_{\pi_e(i)}$, $\forall i \in \Sigma_A$ where $x \in \{-1, 1\}^{|\Sigma_B|}$. It is worth keeping in mind that $x \circ \pi_e \in \{-1, 1\}^{|\Sigma_A|}$.

- (1) Sample an edge $e = (a, b) \in E$. Let π_e be the associated projection constraint.
- (2) Sample $x \sim \{-1, 1\}^{|\Sigma_B|}$, $y \sim \{-1, 1\}^{|\Sigma_A|}$
- (3) Sample $z \sim N_{1-2\delta}((x \circ \pi_e)y)$
- (4) Accept iff $f_a(z) = f_b(x)f_a(y)$

First, note that if the labels assigned to vertices a, b satisfied π_e and f_a, f_b are valid long codes, then $f_a(x \circ \pi_e) = f_b(x)$ for all $x \in \{-1, 1\}^{|\Sigma_B|}$. Second, Step 4 of the verifier can be expressed as a three variable linear equation:

$$(4.8) \quad f_a(z)f_b(x)f_a(y) = 1$$

Proposition 4.5. (Completeness) If $\text{Gap}_\delta \text{LC}$ instance $((A \sqcup B, E), \Sigma_A, \Sigma_B, \{\pi_e\}_{e \in E})$ satisfies all constraints, then there exists code words f_a, f_b such that

$$(4.9) \quad \mathbb{P}[\text{Consistency Test Accepts}] \geq 1 - \delta$$

Proof. Suppose the assignment satisfies all constraints of the $\text{Gap}_\epsilon \text{LC}$ instance above. Set $f_a(x_1, \dots, x_{\Sigma_A}) = x_{\mathfrak{A}(a)}$, $f_b(x_1, \dots, x_{\Sigma_B}) = x_{\mathfrak{B}(b)}$ for some edge $e = (a, b)$. It must hold that:

$$\begin{aligned} \mathbb{P}[\text{Consistency Test Accepts}] &= \mathbb{P}_{x,y,z}[f_a(z) = f_b(x)f_a(y)] \\ &= \mathbb{P}_{x,y,\mu}[(x \circ \pi_e)_{\mathfrak{A}(a)}y_{\mathfrak{A}(a)}\mu_{\mathfrak{A}(a)} = x_{\mathfrak{B}(b)}y_{\mathfrak{A}(a)}] \quad (\mu \sim N_{1-2\delta}(1^{|\Sigma_A|})) \\ &= \mathbb{P}_{x,\mu}[(x \circ \pi_e)_{\mathfrak{A}(a)}\mu_{\mathfrak{A}(a)} = x_{\mathfrak{B}(b)}] \\ &= \mathbb{P}_{x,\mu}[x_{\pi_e(\mathfrak{A}(a))}\mu_{\mathfrak{A}(a)} = x_{\mathfrak{B}(b)}] \\ &= \mathbb{P}_{x,\mu}[x_{\mathfrak{B}(b)}\mu_{\mathfrak{A}(a)} = x_{\mathfrak{B}(b)}] \\ &= 1 - \delta \end{aligned}$$

□

4.5. Folded Functions. A technical detail requires that functions queried f_a, f_b be *folded* in the sense that:

Definition 4.6. A boolean function $f : \{-1, 1\}^n \rightarrow \{-1, 1\}$ is said to be folded if for every $x \in \{-1, 1\}^n$, $f(-x) = -f(x)$.

These functions are also said to be *odd*. Direct calculation shows that *folded* functions have their Fourier weights concentrated on subsets of odd cardinality. In particular, for boolean function f folded: $\widehat{f}(\emptyset) = 0$. Why would we need this property? One issue that arises with the predicates shown in Equation 4.8 is when all functions are simply the constant function $f_a = f_b = 1$. This passes the Noisy BLR test with probability one. This can be avoided by restricting the functions to be folded. Of course, dictator functions are folded. Furthermore, we do not need to increase the number of queries to enforce this requirement: for any two pair of queries $(x, -x)$ and a queried function f , if the test queries $f(x)$, the value $f(-x)$ can be subsequently computed as $-f(x)$. Similarly, if $f(-x)$ is queried, then $f(x)$ is computed as $-f(-x)$. Hence, we lose no generality by assuming the functions we query are folded.

4.6. Soundness and Decoding a Labeling.

Proposition 4.7. Suppose that for a $\text{Gap}_{\epsilon}\text{LC}$ instance, the consistency test above accepts with at least $\frac{1}{2} + \delta$, then we wish to prove there exists a labeling which satisfies at least δ of the projection constraints.

Note that the soundness criterion intuitively implies the use of a method to extract out a labeling from a high enough acceptance probability. This is the “local-correctability” property of linear functions coming into play. We first prove the following lemma:

Definition 4.8. For $S \subseteq \Sigma_A$ and $\pi : \Sigma_A \rightarrow \Sigma_B$, define $\pi_2(S) = \{j \in \Sigma_B \mid |\pi^{-1}(j) \cap S| \text{ is odd}\}$

Lemma 4.9. Assume the test accepts with at least $\frac{1}{2} + \delta$ probability, then for at least δ fraction of the edges (a, b) , the associated codes $f_a : \{-1, 1\}^{|\Sigma_A|} \rightarrow \{-1, 1\}$, $f_b : \{-1, 1\}^{|\Sigma_B|} \rightarrow \{-1, 1\}$ satisfy the following inequality:

$$(4.10) \quad \sum_{S \subseteq [n]} \widehat{f}_a(S)^2 \widehat{f}_b(\pi_2(S)) (1 - 2\delta)^{|S|} \geq \delta$$

Proof. The proof is remarkably similar to that of Lemma 4.4. From our assumption, it is immediate that:

$$(4.11) \quad \frac{1}{2} + \frac{1}{2} \mathbb{E}_{(u,v), x, y, \mu} [f_b(x) f_a(y) f_a((x \circ \pi) y \mu)] \geq \frac{1}{2} + \delta \quad (\mu \sim N_{1-2\delta}(1^{|\Sigma_A|}))$$

which implies that:

$$(4.12) \quad \mathbb{E}_{(u,v), x, y, \mu} [f_b(x) f_a(y) f_a((x \circ \pi) y \mu)] \geq 2\delta$$

By an averaging argument, there must exist at least δ fraction of the edges which satisfy:

$$(4.13) \quad \mathbb{E}_{x, y, \mu} [f_b(x) f_a(y) f_a((x \circ \pi) y \mu)] \geq \delta$$

Hence, through a Fourier-analytic argument:

$$\begin{aligned}
\mathbb{E}_{x,y,\mu} [f_b(x)f_a(y)f_a((x \circ \pi)y\mu)] &\geq \delta \implies \\
\mathbb{E}_{x,y} [f_b(x)f_a(y)\mathbb{E}_\mu [f_a((x \circ \pi)y\mu)]] &\geq \delta \implies \\
\mathbb{E}_{x,y} [f_b(x)f_a(y)\mathcal{T}_{1-2\delta}(f_a)(x \circ \pi)y)] &\geq \delta \implies \\
\mathbb{E}_x [f_b(x)(f * \mathcal{T}_{1-2\delta}(f_a))(x \circ \pi)] &\geq \delta \implies \\
\sum_S \widehat{f_b}(S)\widehat{f_a}(\pi^{-1}(S))^2(1-2\delta)^{|S|} &\geq \delta \implies \\
\sum_S \widehat{f_b}(\pi(S))\widehat{f_a}(S)^2(1-2\delta)^{|S|} &\geq \delta \implies \\
\sum_S \widehat{f_b}(\pi_2(S))\widehat{f_a}(S)^2(1-2\delta)^{|S|} &\geq \delta
\end{aligned}$$

The last implication follows from the observation that:

$$\chi_S(x \circ \pi) = \prod_{i \in S} x_{\pi(i)} = \prod_{i \in \pi(S)} x_i = \prod_{i \in \pi_2(S)} x_i = \chi_{\pi_2(S)}(x)$$

□

We can now finish the soundness proof by decoding the labeling from our supplied proof string π . The actual procedure is not complicated:

- (1) For each $a \in A$, do the following steps:
- (2) Sample $S \sim (\widehat{f_a}(S))^2$
- (3) Set $\mathfrak{A}(a) \stackrel{R}{\leftarrow} S$ (uniformly sample from S)

Do the same for each $b \in B$ except sample the subset $T \sim \widehat{f_b}^2(T)$ from the Fourier spectrum of f_b :

- (1) For each $b \in B$, do the following steps:
- (2) Sample $T \sim (\widehat{f_b}(T))^2$
- (3) Set $\mathfrak{B}(b) \stackrel{R}{\leftarrow} T$ (uniformly sample from T)

Since f_a, f_b are assumed to be folded (odd) functions, both $\widehat{f_a}(\emptyset) = \widehat{f_b}(\emptyset) = 0$, so we will never sample the empty set in any of the procedures above. Remarkably, the procedure above claims that finding a labeling which satisfies an edge with sufficiently high probability amounts to just sample from the Fourier spectrums of the functions f_a, f_b ! We now aim to show the inequality:

Lemma 4.10.

$$(4.14) \quad \mathbb{P}_{(a,b) \in E}[(a,b) \text{ is satisfied}] \geq \sum_S \sum_{T \subseteq \pi(S)} \widehat{f_a}(S)^2 \widehat{f_b}(T)^2 \cdot \frac{1}{|S|}$$

Proof. Here is one way to get a label which satisfies a chosen edge (a,b) . The first procedure samples a subset from the Fourier spectrum of f_a . The second procedure then samples from the Fourier spectrum of f_b and receives a subset $T \subseteq \pi(S)$. The second procedure then uniformly samples from T to receive some label $t \in T$. Since T is contained in the image of S under π , there must exist at least one element in $s \in S$ such that $\pi(s) = t$, showing that this event occurs with probability at least $\frac{1}{|S|}$ □

Proposition 4.11. (Soundness) For the labeling scheme constructed above, the probability:

$$(4.15) \quad \mathbb{P}_{(a,b) \in E}[(a,b) \text{ is satisfied}] \geq 4\delta^4$$

Proof. By Lemma 4.10:

$$\begin{aligned}
\mathbb{P}_{(a,b) \in E}[(a,b) \text{ is satisfied}] &\geq \sum_S \sum_{T \subseteq \pi(S)} \hat{f}_a(S)^2 \hat{f}_b(T)^2 \cdot \frac{1}{|S|} \\
&\geq \sum_S \hat{f}_a(S)^2 \hat{f}_b(\pi_2(S))^2 \frac{1}{|S|} \\
&= \sum_S \left(\hat{f}_a(S) \hat{f}_b(\pi_2(S)) \frac{1}{\sqrt{|S|}} \right)^2 \cdot \left(\sum_S \hat{f}(S)^2 \right) \\
&\geq \left(\sum_S \hat{f}_a(S)^2 \hat{f}_b(\pi_2(S)) \frac{1}{\sqrt{|S|}} \right)^2 \quad (\text{Cauchy-Schwarz}) \\
&\geq 4\delta \left(\sum_S \hat{f}_a(S)^2 \hat{f}_b(\pi_2(S)) (1 - 2\delta)^{|S|} \right)^2
\end{aligned}$$

The last inequality holds since $\frac{1}{\sqrt{x}} \geq 2\sqrt{\delta}(1 - 2\delta)^x$ for all $x, \delta > 0$. By virtue of our bound derived in Lemma 4.9, we find that for at least an δ of the edges, $\sum_S \hat{f}_a(S)^2 \hat{f}_b(\pi_2(S)) (1 - 2\delta)^{|S|} \geq \delta$. Hence:

$$(4.16) \quad \mathbb{P}_{(a,b) \in E}[(a,b) \text{ is satisfied}] \geq 4\delta \cdot \delta^3 = 4\delta^4$$

□

By Theorem 3.5, we know that $\text{Gap}_{4\delta^4}\text{LC}$ is NP-hard for every $\delta \in [0, \frac{1}{2}]$. Since $\delta \geq 4\delta^4$ for $\delta \in [0, \frac{1}{2}]$. The argument for the completeness of this PCP verifier still holds if we set our δ to be $\delta' = 4\delta^4$. This demonstrates that we gave a PCP system for the NP-hard $\text{Gap}_{\delta'}\text{LC}$:

Proposition 4.12. Suppose for an input $\text{Gap}_{\delta}\text{LC}$ instance say \mathcal{I} and associated proof string π causes the test to pass with at least $\frac{1}{2} + \delta$:

$$(4.17) \quad \mathbb{P}[\text{Consistency Test Passes}] \geq \frac{1}{2} + \delta$$

then there exists a assignment for \mathcal{I} such that it satisfies at least $\delta' = 4\delta^4$ fraction of edges in \mathcal{I} .

This gives a PCP system for NP as desired.

5. UNIQUE GAMES

5.1. Definitions. The PCP Theorem culminated in a proof of the NP-hardness of Label Cover by Håstad. Although these results gave proofs of the NP-hardness of $\text{Gap}_{\frac{7}{8}+\epsilon, 1-\epsilon}\text{-Max3SAT}$ and $\text{Gap}_{\frac{1}{2}+\epsilon, 1-\epsilon}\text{-MaxE3Lin}$, similar hardness proofs for other canonical problems such as **MaxCut** didn't seem to follow from these ideas. In his seminal paper, Khot proposed a relaxation of the Label Cover Problem [11]. The instances of this relaxed version are called *Unique Games*:

Definition 5.1. A *Unique Label Cover Problem* with m labels ($\text{UniqueLC}(m)$) instance \mathcal{U} is defined by a bipartite graph $(A \sqcup B, E)$ where $|A| = |B| = n$ for some $n \in \mathbb{N}$, finite alphabet $\Sigma_A = \Sigma_B = \Sigma$ such that $|\Sigma| = m$, and a set of *permutations* $\pi_e : [m] \rightarrow [m]$ for every edge $e \in E$. Define an *assignment* as consisting of a map $\sigma : A \sqcup B \rightarrow [m]$. An edge $e = (a, b) \in E$ is said to be satisfied by this assignment if the assignment is compatible with projection π_e :

$$(5.1) \quad \pi_e(\sigma(a)) = \sigma(b)$$

The value of this game will be

$$(5.2) \quad \text{Opt}(\mathcal{G}) = \max_{\sigma} \mathbb{E}_{e \sim E} [e \text{ satisfied}]$$

In other words, the value will be the largest fraction of edges satisfied by any assignment to the vertices. The corresponding gap problem for Label Cover, $\text{Gap}_{\alpha, \beta} \text{UniqueLC}(m)$, is defined as the promise problem:

$$\begin{aligned} \text{YES} &= \{\mathcal{U} \mid \text{Opt}(\mathcal{U}) \geq \beta\} \\ \text{NO} &= \{\mathcal{U} \mid \text{Opt}(\mathcal{U}) < \alpha\} \end{aligned}$$

In the case of perfect completeness, we abbreviate $\text{Gap}_{\alpha, 1} \text{UniqueLC}(m)$ as simply $\text{Gap}_{\alpha} \text{UniqueLC}(m)$.

In addition, Khot formulated the *Unique Games Conjecture* and utilized it to prove several inapproximability results assuming the conjecture is true.

Conjecture 5.2. (Unique Games Conjecture [11]) For any constant $\delta > 0$, there exists sufficiently large $m \in \mathbb{N}$ such that $\text{Gap}_{\delta, 1-\delta} \text{UniqueLC}(m)$ is NP-hard.

Remark. The $1 - \delta$ constant is crucial to the validity of the conjecture. For an instance of $\text{UniqueLC}(m)$ for any m with a *guaranteed* solution, there exists a polynomial-time algorithm finding an assignment which satisfies all projection constraints: First, start with a vertex and set it to a label. If the vertex is an endpoint of an edge e , follow e to the other side and find a label which satisfies as many neighbors as possible. Repeat in a breadth-first search fashion. In virtue of the projection maps being permutations, this amounts to searching for the guaranteed solution in time $O(mn^2)$.

Example 5.3. The MaxCut problem for an input graph $G = (V, E)$ can be cast as a $\text{UniqueLC}(|V|)$ instance. The two partitions of the bipartite graph A, B will be indexed by the vertices V . Draw an edge between two vertices v_1, v_2 if $(v_1, v_2) \in E$. Set the alphabet to be $\Sigma = \{-1, 1\}$ and the projection maps to be the “swap” map $-1 \mapsto 1, 1 \mapsto -1$.

Example 5.4. MaxE2LinMod_p for prime p denotes the problem of finding an assignment which maximizes the number of satisfied linear constraints consisting of exactly two variables over field \mathbb{F}_p . An example of an instance of this problem over variables x_1, x_2, x_3, x_4 is shown below:

$$\begin{aligned} x_1 + x_3 &= 3 \\ x_2 + x_4 &= 2 \\ x_1 + x_4 &= 1 \end{aligned}$$

An instance of this problem can also be translated as an instance of $\text{UniqueLC}(p)$.

Definition 5.5. A promise problem P is said to be *UG-hard* if there is some constant $\delta > 0$ such that for all $m \in \mathbb{N}$, there exists a polynomial-time reduction f from $\text{Gap}_{\delta, 1-\delta} \text{UniqueLC}(m)$ to P , in the sense that for a $\text{Gap}_{\delta, 1-\delta} \text{UniqueLC}(m)$ instance \mathcal{U} :

$$\begin{aligned} \text{Opt}(\mathcal{U}) \geq 1 - \delta &\implies f(\mathcal{U}) \in \text{Yes} \\ \text{Opt}(\mathcal{U}) < \delta &\implies f(\mathcal{U}) \in \text{No} \end{aligned}$$

6. UG-HARDNESS OF MAXCUT

6.1. Intuitions. Recall that the keen insight behind Håstad's 3-bit PCP for NP is the embedding of a dictatorship test within a Label Cover instance. In a similar spirit, the proof for the optimality of the Goemans-Williamson algorithm will hinge on crafting a clever dictatorship test embedded within a MaxCut instance. For the sake of posterity, we first introduce the Goemans-Williamson algorithm.

6.2. Goemans-Williamson Algorithm for MaxCut. Example 1.3 presented an LP-based approximation algorithm for MaxCut. Goemans and Williamson [6] designed an SDP to give an α_{GW} -approximation algorithm for MaxCut where:

$$(6.1) \quad \alpha_{GW} = \min_{-1 \leq \rho \leq 1} \frac{2 \cos^{-1}(\rho)}{\pi} \approx 0.87856$$

To begin, a semi-definite program is a generalization of a linear program where instead of optimizing over a vector of variables \vec{x} , the program considers a positive semi-definite matrix of variables i.e a matrix whose eigenvalues are non-negative.

Definition 6.1. A *semi-definite program* is an optimization problem of the following form: let $C, A^{(1)}, \dots, A^{(m)} \in \mathbb{R}^{n \times n}$ be $n \times n$ real-valued matrices and $b^{(1)}, \dots, b^{(m)} \in \mathbb{R}$ be scalars. The objective is to find a positive semi-definite matrix $X \in \mathbb{R}^{n \times n}$ such that:

$$\begin{aligned} & \max C \odot X \\ & \text{under constraints } A^{(k)} \odot X \leq b_k \\ & \quad \forall k \in [m] \end{aligned}$$

where \odot denotes the Schur product or entry-wise product:

$$A \odot B = \sum_{1 \leq i, j \leq n} A_{ij} B_{ij} \quad A, B \in \mathbb{R}^{n \times n}$$

An equivalent formulation considers inner products between pairs of real-valued vectors:

$$\begin{aligned} & \max \sum_{i, j} c_{ij} \langle v_i, v_j \rangle \\ & \text{under constraints } a_{ij}^{(k)} \langle v_i, v_j \rangle \leq b^{(k)} \\ & \quad v_1, \dots, v_n \in \mathbb{R}^n, \quad \forall k \in [m] \end{aligned}$$

Note that this form also subsumes quadratic programming by setting $c_{ij} = a_{ij}^k = b_{ij}^k = 0$ for all $i \neq j$ and $k \in [m]$.

The Goemans-Williamson algorithm concerns the solution to the semi-definite program given some undirected graph $G = (V, E)$:

$$\begin{aligned} & \max \sum_{i, j} \frac{1 - \langle v_i, v_j \rangle}{2} \\ & \langle v_i, v_i \rangle = 1 \text{ for all } i \in V \end{aligned}$$

Let us first show that the program is a relaxation of the integer program crafted in Example 1.3. Indeed, if we set any unit vector \vec{u} such that for a cut defined by (S, \bar{S}) :

$$\vec{v}_i = \begin{cases} u & \text{if } i \in S \\ -u & \text{if } i \notin S \end{cases}$$

A direct calculation yields that the term $\frac{1 - \langle v_i, v_j \rangle}{2} = 0$ if $v_i = v_j = u$ else it is equal to 1. Thus, by applying this observation to the maximum cut,

$$\text{Opt}_{GW}(G) \geq \text{MaxCut}(G)$$

So far the semi-definite program appears to be rather simple. The insight made by Goemans and Williamson lies in the rounding procedure of the solution outputted by the program. The procedure proceeds by first drawing a random hyperplane passing through the origin and then taking the partitions of the cut S_+, S_- to be the solutions v_i which lie on positive and negative sides of the hyperplane respectively. We can express the expected size of the cut:

$$\mathbb{E}[|E(S_+, S_-)|] = \sum_{(i,j) \in E} \mathbb{P}[v_i, v_j \text{ lie on different sides of the hyperplane}]$$

Now if θ_{ij} denotes the angle between v_i, v_j , a simple geometric argument shows that:

$$\mathbb{P}[v_i, v_j \text{ lie on different sides of the hyperplane}] = \frac{\theta_{ij}}{\pi}$$

By definition of the dot product:

$$\frac{1 - \langle v_i, v_j \rangle}{2} = \frac{1 - \cos(\theta_{ij})}{2}$$

and by virtue of the magical inequality:

$$\frac{\theta_{ij}}{\pi} \geq \alpha_{GW} \cdot \frac{1 - \cos(\theta_{ij})}{2}$$

by taking the sum over all edges over both sides, the two can be combined to yield that:

$$(6.2) \quad \frac{\mathbb{E}[|E(S_+, S_-)|]}{\text{Opt}_{GW}(f)} \geq \alpha_{GW}$$

Furthermore, the first inequality above implies that the Goemans-Williamson algorithm is indeed an α_{GW} approximation algorithm:

$$(6.3) \quad \frac{\mathbb{E}[|E(S_+, S_-)|]}{\text{MaxCut}(G)} \geq \alpha_{GW}$$

6.3. MaxCut is UG-hard. After introducing the basic background, we are finally ready to show a non-trivial inapproximability result assuming the UGC:

Theorem 6.2. (MaxCut is UG-hard) *The problem $\text{Gap}_{\frac{\cos^{-1}(\rho)}{\pi} + \epsilon, \frac{1-\rho}{2} - \epsilon} \text{MaxCut}$ is UG-hard*

Note that an immediate corollary of Theorem and the UGC is that the Goemans-Williamson algorithm is tight:

Corollary 6.2.1. *Assuming the UGC, if an α_{GW} -approximation algorithm for MaxCut exists, then $P = NP$.*

Proof. Using Theorem 1.9, we see that for small ϵ :

$$\frac{\frac{\cos^{-1}(\rho)}{\pi} + \epsilon}{\frac{1-\rho}{2} - \epsilon} \approx \frac{2 \cos^{-1}(\rho)}{\pi (1-\rho)}$$

Setting $\rho \approx -0.6934$ yields the desired result. \square

Onwards to the proof of Theorem 6.2. We first reason about the relationship between $\text{UniqueLC}(m)$ instances and MaxCut instances. In particular, we wish to demonstrate that proving UG-hardness for MaxCut is equivalent to constructing a 2-query PCP for $\text{GapUniqueLC}(m)$ for all m .

Theorem 6.3. $\text{Gap}_{\alpha,\beta}\text{MaxCut}$ is UG-hard iff there exists a constant $\delta > 0$ such that for all m there exists a 2-query PCP for $\text{Gap}_{\delta,1-\delta}\text{UniqueLC}(m)$ with completeness β and soundness α .

Proof. First assume that for some $\delta > 0$ and all m , there exists a polynomial-time reduction from $\text{Gap}_{\delta,1-\delta}\text{UniqueLC}(m)$ to $\text{Gap}_{\alpha,\beta}\text{MaxCut}$. Let $G = (V, E), \mathcal{C}$ be the graph outputted by reduction and the cut outputted by the cut approximation algorithm. Construct a PCP samples two vertices from $v_1, v_2 \sim V$ by querying the proof tape encoding G, \mathcal{C} and outputs “accept” if $(v_1, v_2) \in \mathcal{C}$. The completeness and soundness parameters immediately follow from definition of $\text{Gap}_{\alpha,\beta}\text{MaxCut}$. Conversely, assume the existence of a 2-query PCP for $\text{Gap}_{\delta,1-\delta}\text{UniqueLC}(m)$ with the above properties. We will calculate the acceptance probabilities of this PCP given a proof string and an input instance by finding the max cut. Let the vertices be the proof locations of the input proof string π . It suffices to draw an edge between two vertices weighted by the probability their corresponding proof locations are queried and accepted by the verifier. \square

Thus, we can proceed by constructing a 2-query PCP for $\text{Gap}_{\delta,1-\delta}\text{UniqueLC}(m)$ with the completeness $\frac{1-\rho}{2} - \epsilon$ and soundness $\frac{\cos^{-1}(\rho)}{\pi} + \epsilon$. As with Håstad’s 3-bit PCP, we aim to construct a dictatorship test which generates the parameters needed. Once again, an assignment of labels to the vertices of a $\text{UniqueLC}(m)$ instance would be encoded into a truth table of a dictator i.e into a *long code*.

6.4. Dictatorship Testing and Majority is the Stablest (MIS). A key difference between the dictatorship test crafted in Håstad’s 3-bit PCP for NP and the one we aim to devise here is the number of queries the verifier can make. The verifier here can only use *two* queries to the proof rather than the three queries afforded to Håstad’s verifier. The dictatorship test we will focus on for this reduction utilizes only two queries to the functions: given input boolean function $f : \{-1, 1\}^n \rightarrow \{-1, 1\}$ and $\rho \in (-1, 0)$:

- (1) Sample $x \sim \{-1, 1\}^n$
- (2) Sample $y \sim N_\rho(x)$
- (3) Accept iff $f(x) \neq f(y)$

Before we analyze the acceptance probability of our dictatorship test, we require a tool bounding the noise sensitivity of a boolean function with small low-degree influence. This is captured by the “Majority is the Stablest” theorem:

Theorem 6.4. (“Majority is the Stablest” (MIS) [14]) *For every $\epsilon > 0$, $\rho \in (-1, 0)$, there exists $\tau > 0$ such that if for all $i \in [n]$, $\text{Inf}_i(f) < \tau$ for function $f : \{-1, 1\}^n \rightarrow [-1, 1]$, then*

$$(6.4) \quad \text{NS}_\rho(f) < \frac{\cos^{-1}(\rho)}{\pi} + \epsilon$$

Recall that the noise sensitivity of a boolean function $f : \{-1, 1\}^n \rightarrow \{-1, 1\}$ is defined as:

$$(6.5) \quad \text{NS}_\rho(f) = \mathbb{P}_{y \sim N_\rho(x), x}[f(x) \neq f(y)]$$

where $y \sim N_\rho(x)$ refers to sampling a string y under the procedure:

$$y_i = \begin{cases} x_i & \text{with probability } \frac{1+\rho}{2} \\ -x_i & \text{with probability } \frac{1-\rho}{2} \end{cases}$$

Now equation 6.5 can be re-expressed in terms of the noise stability of f :

$$(6.6) \quad \text{NS}_\rho(f) = \frac{1}{2} - \frac{1}{2} \text{Stab}_\rho(f)$$

where

$$\text{Stab}_\rho(f) = \mathbb{E}_{y \sim N_\rho(x), x}[f(x)f(y)]$$

Through Fourier-analytic techniques, we derive that:

$$(6.7) \quad \text{NS}_\rho(f) = \frac{1}{2} - \frac{1}{2} \sum_{S \subseteq [n]} \hat{f}(S)^2 \rho^{|S|}$$

To see the significance of the MIS Theorem, we compute the acceptance rate of the dictatorship test above:

$$\begin{aligned} \mathbb{P}[\text{Test Accepts}] &= \mathbb{E}_{x,y} \left[\frac{1}{2} - \frac{1}{2} f(x)f(y) \right] \\ &= \frac{1}{2} - \frac{1}{2} \mathbb{E}_{x,y} [f(x)f(y)] \\ &= \frac{1}{2} - \frac{1}{2} \text{Stab}_\rho(f) \end{aligned}$$

We can now reason about acceptance probability in the context of several classes of functions:

- Dictators pass the test with probability $\frac{1-\rho}{2}$
- The parity functions χ_S , $S \subseteq [n]$ pass with probability $\frac{1-\rho^{|S|}}{2}$. As $|S|$ grows, the acceptance probability tends to $\frac{1}{2}$.
- Majority functions passes the test with probability approaching in the limit $\frac{\cos^{-1}(\rho)}{\pi}$.

What the MIS Theorem roughly tells us is that functions with no “notable” coordinates, i.e all coordinates have small influence, can pass the test with probability at most $\frac{\cos^{-1}(\rho)}{\pi}$. The proof of the soundness of our verifier will exploit the observation that functions with a sufficiently large i^{th} influence for some coordinate i will be “close enough” to a dictator. This property will be used to extract a labeling to the $\text{UniqueLC}(m)$ instance satisfying a large fraction of constraints.

For the analysis of constructed PCP verifier, a generalization of the MIS theorem will serve our purposes:

Theorem 6.5. (*Generalized MIS* [12],[14]) For all $\epsilon > 0$, $\rho \in (0, 1)$, there exists some $\tau > 0$ and finite d such that if $f : \{-1, 1\}^n \rightarrow [-1, 1]$ and for all $i \in [n]$, $\text{Inf}_i^{\leq d}(f) \leq \tau$:

$$(6.8) \quad \frac{1}{2} - \frac{1}{2} \sum_{S \subseteq [n]} \hat{f}(S)^2 \rho^{|S|} < \frac{\cos^{-1}(\rho)}{\pi} + \epsilon$$

where $\text{Inf}_i^{\leq d}(f)$ refers to the i^{th} influence of f restricted to degrees of at most d :

$$(6.9) \quad \text{Inf}_i^{\leq d}(f) = \sum_{\substack{S \subseteq [n], i \in S \\ |S| \leq d}} |S| \cdot \hat{f}(S)^2$$

Note that the image of our function f is contained in the closed interval $[-1, 1]$ rather than the finite set $\{-1, 1\}$ as stated in the ordinary MIS theorem.

6.5. The 2-query PCP. With the MIS theorem, we can begin our analysis of the promised 2-query PCP. Before the procedure, define $(x \circ \pi)_i = x_{\pi(i)}$ for $x \in \{-1, 1\}^n$, $i \in [n]$, $\pi \in S_n$.

- (1) Sample a vertex $a \in A$ uniformly.
- (2) Sample two of its neighbors $b, b' \in B$ uniformly. Let $\pi_{b,a}$ and $\pi_{b',a}$ denote the *inverses* of the constraints associated to $(a, b), (a, b')$ respectively.
- (3) Sample $x \sim \{-1, 1\}^m$ and $y \sim N_\rho(x)$
- (4) Accept if $f_b(x \circ \pi_{b,a}) \neq f_{b'}(y \circ \pi_{b',a})$

For the proof of Theorem 6.2, invoke Theorem 6.5 for $\frac{\epsilon}{2}$ and ρ assumed. This will yield a τ and a degree upper bound d . Set parameter

$$\delta = \frac{\epsilon \tau^2}{8d}$$

6.5.1. Completeness. Suppose we have a **UniqueLC**(m) instance \mathcal{U} such that $\text{Opt}(\mathcal{U}) \geq 1 - \delta$. Through a simple union bound argument, the probability that both $\pi_{b,a}, \pi_{b',a}$ are satisfied by the assignment is at least $1 - 2\delta$. Now if both are indeed satisfied and the test accepts, it must be true that:

$$\begin{aligned} f_b(x \circ \pi_{b,a}) \neq f_{b'}(y \circ \pi_{b',a}) &\iff (x \circ \pi_{b,a})_{\sigma(b)} \neq (y \circ \pi_{b',a})_{\sigma(b')} \\ &\iff x_{\pi_{b,a}(\sigma(b))} \neq y_{\pi_{b',a}(\sigma(b'))} \\ &\iff x_{\sigma(a)} \neq y_{\sigma(a)} \end{aligned}$$

where for the last equivalence, we invoked the assumption that σ satisfies both $\pi_{b,a}, \pi_{b',a}$. Recall that σ is the assignment of labels to the vertices of the bipartite graph. The last expression occurs with probability $\frac{1-\rho}{2}$ by step three of the verification algorithm. Hence,

$$\mathbb{P}[\text{Test accepts}] \geq \frac{(1 - 2\delta)(1 - \rho)}{2}$$

By observing that $\delta < \frac{\epsilon}{2}$, the inequality above reduces to:

$$(6.10) \quad \mathbb{P}[\text{Test accepts}] \geq \frac{(1 - \epsilon)(1 - \rho)}{2} \geq \frac{1 - \rho}{2} - \epsilon$$

as desired.

6.5.2. *Soundness.* To show soundness, we prove the contrapositive, namely if

$$\mathbb{P}[\text{Test accepts}] \geq \frac{\cos^{-1}(\rho)}{\pi} + \epsilon$$

then there exists a labeling which satisfies more than a δ fraction of constraints.

Proof. First:

$$(6.11) \quad \mathbb{P}[\text{Test accepts}] = \mathbb{E}_{a,b,b'} \left[\mathbb{E}_{x,y \sim N_\rho(x)} \left[\frac{1}{2} - \frac{1}{2} f_b(x \circ \pi_{b,a}) f_{b'}(y \circ \pi_{b',a}) \right] \right] \geq \frac{\cos^{-1}(\rho)}{\pi} + \epsilon$$

An averaging argument on the vertex $a \in A$ tells us that, since the test passes with at least $\frac{\cos^{-1}(\rho)}{\pi} + \epsilon$ probability, there must exist at least $\epsilon/2$ fraction of the vertices in A such that the test conditioned on picking a from this fraction passes with probability at least $\frac{\cos^{-1}(\rho)}{\pi} + \epsilon/2$. Otherwise, if there existed fewer than a $\epsilon/2$ fraction such that the test passed with at least this probability, then the total probability of the test passing is *at most* $(\epsilon/2) \cdot 1 + (1 - \epsilon/2)(\frac{\cos^{-1}(\rho)}{\pi} + \epsilon/2) < \frac{\cos^{-1}(\rho)}{\pi} + \epsilon$, which is a contradiction. Let us label the vertices picked from this $\epsilon/2$ fraction as *good* vertices.

So say we picked one of these good vertices say a . Let us define the below function:

Definition 6.6. Define $g_a : \{-1, 1\}^m \rightarrow [-1, 1]$ as

$$(6.12) \quad g_a(x) = \mathbb{E}_b [f_b(x \circ \pi_{b,a})]$$

where the expectation is drawn uniformly over the neighbors b of a .

This allows us to re-express the inequality 6.11:

$$(6.13) \quad \mathbb{E}_{b,b'} \left[\mathbb{E}_{x,y \sim N_\rho(x)} \left[\frac{1}{2} - \frac{1}{2} f_b(x \circ \pi_{b,a}) f_{b'}(y \circ \pi_{b',a}) \right] \right]$$

$$(6.14) \quad = \mathbb{E}_{x,y \sim N_\rho(x)} \left[\frac{1}{2} - \frac{1}{2} \mathbb{E}_b [f_b(x \circ \pi_{b,a})] \mathbb{E}_{b'} [f_{b'}(y \circ \pi_{b',a})] \right]$$

$$(6.15) \quad = \frac{1}{2} - \frac{1}{2} \mathbb{E}_{x,y \sim N_\rho(x)} [g_a(x) g_a(y)]$$

$$(6.16) \quad = \frac{1}{2} - \frac{1}{2} \text{Stab}_\rho(g_a)$$

$$(6.17) \quad \geq \frac{\cos^{-1}(\rho)}{\pi} + \epsilon/2$$

By invoking the contrapositive of the generalized MIS theorem (Theorem 6.5) on g_a , we deduce that there must exist some index $i_a \in [m]$ such that:

$$(6.18) \quad \inf_{i_a}^{\leq d}(g_a) \geq \tau$$

Fortunately, there is a method to equate the Fourier coefficients g_a with those of f_b :

$$\begin{aligned}
g_a(x) &= \mathbb{E}_b [f_b(x \circ \pi_{b,a})] \\
&= \mathbb{E}_b \left[\sum_{S \subseteq [n]} \hat{f}_b(S) \chi_S(x \circ \pi_{b,a}) \right] \\
&= \mathbb{E}_b \left[\sum_{S \subseteq [n]} \hat{f}_b(S) \chi_{\pi_{b,a}(S)}(x) \right] \\
&= \mathbb{E}_b \left[\sum_{S \subseteq [n]} \hat{f}_b(S) \chi_{\pi_{b,a}^{-1}(S)}(x) \right] \\
&= \sum_{S \subseteq [n]} \mathbb{E}_b [\hat{f}_b(\pi_{b,a}^{-1}(S))] \chi_S(x)
\end{aligned}$$

where the last equality follows from reparameterizing the sum. By expanding g_a on the left-hand side of the equality into its Fourier expansion, from the orthogonality of characters:

$$(6.19) \quad \hat{g}_a(S) = \mathbb{E}_b [\hat{f}_b(\pi_{b,a}^{-1}(S))]$$

Starting from inequality 6.18,

$$\begin{aligned}
\tau &\leq \inf_{i_a}^{\leq d} (g_a) = \sum_{\substack{S \subseteq [n], i_a \in S \\ |S| \leq d}} \hat{g}_a^2(S) \\
&= \sum_{\substack{S \subseteq [n], i_a \in S \\ |S| \leq d}} \left(\mathbb{E}_b [\hat{f}_b(\pi_{b,a}^{-1}(S))] \right)^2 \\
&\leq \sum_{\substack{S \subseteq [n], i_a \in S \\ |S| \leq d}} \mathbb{E}_b [\hat{f}_b^2(\pi_{b,a}^{-1}(S))] \\
&= \mathbb{E}_b \left[\sum_{\substack{S \subseteq [n], i_a \in S \\ |S| \leq d}} \hat{f}_b^2(\pi_{b,a}^{-1}(S)) \right] \\
&= \mathbb{E}_b \left[\inf_{\pi_{b,a}^{-1}(i_a)}^{\leq d} (f_b) \right]
\end{aligned}$$

The inequality uses Cauchy-Schwarz. We use another averaging argument to see that there must exist at least a $\tau/2$ fraction of a 's neighbors b such that

$$\inf_{\pi_{b,a}^{-1}(i_a)}^{\leq d} (f_b) \geq \tau/2$$

otherwise the total influence term would be at most $\tau/2 \cdot 1 + (1 - \tau/2)(\tau/2) < \tau$. For each b in that $\tau/2$ fraction, we pick a label uniformly at random from the set:

$$S_b = \{\ell \mid \inf_{\ell}^{\leq d} (f_b) \geq \tau/2\}$$

which must be *non-empty* by the averaging argument made above. Notice that the label picked will satisfy $\pi_{b,a}$ by construction. We can thus lower bound the probability of constraint $\pi_{b,a}$ being satisfied:

$$(6.20) \quad \mathbb{P}[\pi_{(b,a)} \text{ is satisfied}] \geq \frac{\epsilon}{2} \frac{\tau}{2} \frac{1}{|S_b|}$$

Through a Fourier-analytic argument, we can upper-bound $|S_b|$:

$$\frac{|S_b|\tau}{2} \leq \sum_{i=1}^{|S_b|} \text{Inf}_i^{\leq d}(f_b) \leq \sum_{i=1}^m \text{Inf}_i^{\leq d}(f_b) = \sum_{\substack{S \subseteq [m] \\ |S| \leq d}} |S| \hat{f}^2(S) \leq d \sum_{S \subseteq [m]} \hat{f}^2(S) = d$$

This yields that $|S_b| \leq \frac{2d}{\tau}$. So the probability bound of 6.20 would become:

$$(6.21) \quad \mathbb{P}[\pi_{(b,a)} \text{ is satisfied}] \geq \frac{\epsilon}{2} \frac{\tau}{2} \frac{\tau}{2d} = \frac{\epsilon \tau^2}{8d} = \delta$$

as desired. This completes the proof of Theorem 6.2. \square

REFERENCES

1. Sanjeev Arora and Boaz Barak, *Computational complexity: a modern approach*, Cambridge University Press, 2009.
2. Sanjeev Arora, Carsten Lund, Rajeev Motwani, Madhu Sudan, and Mario Szegedy, *Proof verification and the hardness of approximation problems*, Journal of the ACM (JACM) **45** (1998), no. 3, 501–555.
3. Sanjeev Arora and Shmuel Safra, *Probabilistic checking of proofs: A new characterization of np*, Journal of the ACM (JACM) **45** (1998), no. 1, 70–122.
4. Mihir Bellare, Oded Goldreich, and Madhu Sudan, *Free bits, pcps, and nonapproximability—towards tight results*, SIAM Journal on Computing **27** (1998), no. 3, 804–915.
5. Irit Dinur, *The pcp theorem by gap amplification*, Journal of the ACM (JACM) **54** (2007), no. 3, 12–es.
6. Michel X Goemans and David P Williamson, *Improved approximation algorithms for maximum cut and satisfiability problems using semidefinite programming*, Journal of the ACM (JACM) **42** (1995), no. 6, 1115–1145.
7. Venkatesan Guruswami and Ryan O’Donnell, *Cse 533: The pcp theorem and hardness of approximation*, 2005.
8. Prahladh Harsha, *Limits of approximation algorithms: Pcps and unique games (tifr/imsc lecture notes, spring 2010)*, (2010).
9. Prahladh Harsha, Moses Charikar, Matthew Andrews, Sanjeev Arora, Subhash Khot, Dana Moshkovitz, Lisa Zhang, Ashkan Aazami, Dev Desai, Igor Gorodezky, et al., *Limits of approximation algorithms: Pcps and unique games (dimacs tutorial lecture notes)*, arXiv preprint arXiv:1002.3864 (2010).
10. Johan Håstad, *Some optimal inapproximability results*, Journal of the ACM (JACM) **48** (2001), no. 4, 798–859.
11. Subhash Khot, *On the power of unique 2-prover 1-round games*, Proceedings of the thirty-fourth annual ACM symposium on Theory of computing, 2002, pp. 767–775.
12. Subhash Khot, Guy Kindler, Elchanan Mossel, and Ryan O’Donnell, *Optimal inapproximability results for max-cut and other 2-variable csps?*, SIAM Journal on Computing **37** (2007), no. 1, 319–357.
13. Subhash Khot and Nisheeth K Vishnoi, *On the unique games conjecture*, FOCS, vol. 5, Citeseer, 2005, p. 3.
14. Elchanan Mossel, Ryan O’Donnell, and Krzysztof Oleszkiewicz, *Noise stability of functions with low influences: invariance and optimality*, 46th Annual IEEE Symposium on Foundations of Computer Science (FOCS’05), IEEE, 2005, pp. 21–30.
15. Ryan O’Donnell, *Analysis of boolean functions*, Cambridge University Press, 2014.
16. Ran Raz, *A parallel repetition theorem*, SIAM Journal on Computing **27** (1998), no. 3, 763–803.
17. Luca Trevisan, *On khot’s unique games conjecture.*, Bulletin (New Series) of the American Mathematical Society **49** (2012), no. 1.

Email address: ehkim@cs.unc.edu